

NISHANK KOUL

+91 9873446506 koulnishank5@gmail.com nishankkoul nishank-koul nishank

SUMMARY

DevOps Engineer with expertise in automating infrastructure, streamlining deployments, and delivering scalable, resilient cloud-native solutions.

EDUCATION

PES University - Bengaluru, India

Dec 2021 - May 2025

Bachelor of Technology: Computer Science

TECHNICAL SKILLS

Languages: Python, Bash

Cloud Platforms: Amazon Web Services (AWS), Google Cloud Platform (GCP)

CI/CD Tools: Jenkins, GitHub Actions

Containerization: Docker, Kubernetes

Monitoring: Prometheus, Grafana, ELK Stack

Infrastructure as Code: Terraform, Ansible

EXPERIENCE

Stringify AI | *DevOps Engineer*

Feb 2025 - Present

- Designed and containerized cloud-native applications using **multi-stage Docker builds**, and deployed them on **Google Cloud**, resulting in a **30% reduction in image size** and significantly improving application startup time and resource efficiency.
- Configured a **Global HTTP(S) Load Balancer** with **backend services as Network Endpoint Groups (NEGs)** and **CDN cache enabled** to accelerate content delivery.
- Provisioned a **production-grade PostgreSQL database on GCP Compute Engine** with SSH access tightly controlled through a bastion host; implemented **SSH tunneling on PgAdmin4** for secure local connectivity and visualization.
- Streamlined CI/CD processes using GitHub Actions, **accelerating deployment cycles by 40%** while ensuring consistency, reliability, and faster time-to-market.
- Integrated **SonarQube** for **SAST** in the CI/CD pipelines, enabling automated detection of code vulnerabilities, enforcing security best practices, and improving code quality across deployments.
- Enabled Cloud Audit Logging in production environment for **enhanced security and compliance, ensuring comprehensive audit trails** for administrative activities and data access in Google Cloud.
- Created **custom events on Google Tag Manager** and integrated them with Google Analytics, enabling **precise tracking of user interactions and providing actionable insights** to optimize marketing strategies and improve conversion rates.

Bimaplan | *DevOps Engineer Intern*

Sep 2024 - Feb 2025

- Developed Python scripts for AWS Lambda functions to automatically shut down EC2 instances in the Dev and UAT environments during non-business hours, leading to a **25% reduction in overall cloud costs** by optimizing resource utilization and minimizing idle time.
- Orchestrated **zero-touch deployment** by automating 90% of AWS infrastructure provisioning using Terraform; enabled Disaster Recovery through cross-region replication to ensure business continuity.
- Executed the setup of a **read replica for the RDS Database** to enhance availability and scalability, **improving read query performance by 40%** and reducing downtime risks.
- Refined Jenkins CI/CD pipelines** across Dev, UAT, and Prod by integrating Terraform, ensuring **100% consistency** in provisioning. **Established backup strategies** for pipeline code and statefiles, reducing rollback time by **60%**.
- Delivered an efficient **API Gateway rate-limiting strategy** based on traffic analysis to enhance performance and prevent abuse; integrated **CloudWatch alarms** with Slack for real-time alerting on HTTP 429 errors.

PROJECTS

Celestia Validator Node Deployment on Mocha-4 Testnet | *Blockchain, Ansible, AWS EC2, Prometheus, Grafana* |

- Built an end-to-end Ansible playbook to **automate Celestia validator node provisioning, reducing manual setup time by 80%** and ensuring consistent deployments with zero configuration drift.
- Configured a **Grafana-based monitoring system with custom dashboards** to track node performance metrics, including block height, sync status, and resource utilization in real-time, enhancing operational visibility and reducing incident resolution time by **50%**.
- Developed industry-standard security protocols by applying encryption and access **restrictions for sensitive credentials using Ansible Vault** and designed rollback mechanisms, **reducing downtime risk by 30%** and improving validator resilience.

Scalable LLM Inference Service with Ollama | *LLMs, Flask, Docker, AWS EKS, K6.io, GitHub Actions* |

- Engineered a **scalable LLM inference service** with Ollama, integrating the moondream model. Developed a Dockerized Flask API and orchestrated the service on AWS EKS for **high availability and scalability**.
- Accelerated application performance by identifying and resolving memory allocation bottlenecks during **Load Testing with K6.io, improving container accessibility and response times**.
- Executed auto-scaling strategies, increasing the successful request response rate from **53.66% to 85.49%**.

CERTIFICATION

GCP Certified Associate Cloud Engineer

Achieved in May 2025

AWS Certified Cloud Practitioner

Achieved in Oct 2024